

Report – Jerry Liu 20186440

Description of data

A sizeable dataset was desired for this project and the actual application of the B+ tree. Taking inspiration from the example given in the assignment, an online Netflix dataset from the website data.world (<https://data.world/chasewillden/netflix-shows>) was taken for use which describes a set of movies. The total dataset was shrunk to 200 entries in total for this assignment.

The data is a collection of movies and descriptions of their properties as well as user feedback. For example, it has the title, rating and a description of what the rating means, in addition to the release date. It also includes the number of users that rated the movie, the rating scores and rating descriptions for more information on what the people thought of the movie.

Description of B+ tree

The B+ tree follows the standard implementation of a B+ tree, with a few customizations due to the unique structure of the keys and values. In addition to the regular B+ tree characteristics, this B+ tree has keys with two attributes, a primary attribute for comparison and a secondary attribute for comparison. This means that when ordering keys and performing searches, the primary key needs to be compared followed by the secondary key if the primary key values are equal for exact precision.

Furthermore, the other main customization was that most B+ trees seem to have leaf nodes with single key value pairs. To follow the diagrams given in the original assignment, the leaf nodes have no repeated keys, and if there are multiple entries with identical keys (both primary and secondary), then they are put together in an ordered list matching the index of the specified key.

The schema of the B+ tree is quite simple. The database is relatively simple and can be represented as one table with a list of movies with unique tuple ids. This is straightforward and allows for the tree to reference it easily during operations. The constraints were that the number of tuples must be at least 20, and have at least 4 attributes for each tuple, including its own identifier. The tree itself had key value pairs where root and internal nodes had values of pointers to child nodes and leaf nodes had values of tuple ids. The order of the B+ tree sets the range of key value pairs a root or internal/leaf node can have. In this case, the root node must have from 1 to one less than the order, and other nodes must have from the order/2 -1 to order-1 pairs, with the ceiling of the order/2 being taken.

There are 200 entries in total, so the number of tuples that exist in the database is 200, while any number from 0 to 200 of tuples could exist in the tree.

The order of the B+ tree has been set to 3.

The attributes that are used for the key of the B+ tree are the rating description as the primary attribute for the key and the release year for the secondary attribute for the key. These were chosen because of all the attributes, they had the most unique values apart from the user rating score. However, certain tuples had a NULL value for the user rating score which would make it difficult to categorize those entries. Therefore, these two attributes became the most useful ones to use as the key.

Running the program

The program was written using Notepad++ as a .py file. To run the program, any command line editor with Python should be able to be used. For testing during development, an Anaconda 3.6 Python terminal was used, and once reaching the directory of the files, the command “python Bplustree.py” needed to be executed and the program would start automatically

Testing the 6 Operations

All functions were tested while loading tuples between 1 and 10 of the dataset. They all work completely except for delete, which does not have the capability to join or borrow from neighbouring nodes when needed for rebalancing. Instead, if this requirement occurs, the parent node will be deleted, which can propagate up to the root node.

Any time an operation is to be select, the program will ask “What do you want to do: “ and expect a single integer from 1-7 as the input. Following this will be the specific input for each of the 6 operations and exit. All operations were only tested on the netflixEdited.csv file and dataset, because of the customizations that would have to be made for the attribute names and other information for another file that could not be easily stored inside the .csv file itself.

To test load, a provided data file in a .csv format is given. This .csv file can be opened in Microsoft Excel or a similar program to see what the real contents are. In the actual program, once it has run the first input should be a “1” followed by the file name of the data and start and end tids. This means 4 inputs in total after the program has begun running. The data used for testing is included as a file called “netflixEdited.csv”. The tids start at 1, so there is no entry that belongs to an tid of 0, which could lead to errors if referenced.

There is no error check for tids that are out of bounds or invalid, so it is assumed that the start and end are both valid logically and programmatically. The tree will automatically be build. The load function not only reads the entire csv and stores it in a table for retrieval and reference, but it inserts the tids from the start to the end as an initialization of the tree. Due to this, there is no way to start with an empty tree for the testing of the other functions, but an empty tree could be created by deleting all the entries. This also means load must be the first operation that is performed with the program every execution. The print operation can be used to output the current state of the tree and matches the expected output from the loading.

To test print, once data has been loaded and other commands have been inputted, the user should have an expected result for the tree and what it looks like based on their inputs. Testing print simply requires typing “2” and the current state of the tree will be printed in a breadth first manner. Comparison of the output and the expected tree can easily be done by looking at the print result layer by layer. The print function has been tested under all known test cases to work.

To test insert, input “3” as the choice followed by the tid that is desired to be inserted. A print message will be given when the tid is successfully added, or a print message will be given if the tid already exists in the tree. A print operation following the insertion can be used to test that it worked successfully and didn’t cause any problems in the tree. The insert operation

To test delete, input "4" as the choice followed by the tid that is desired to be deleted. If the deletion is successful, a print message will be given that it was successfully deleted, otherwise it will print that it was not successful (most likely due to the tid not existing in the tree). A print operation following the deletion can be used to test that the deletion worked properly and as expected if the tree needed to be rearranged.

To test search, enter "5" as the operation input and it will ask for the value of the primary and secondary key attributes the user would like to use in the search. The program will combine both these attributes into the single key that it is looking for in the tree, and perform a search. It will output the tids that match exactly the primary and secondary key attributes provided by the user. To test this function, a tree with keys that match and don't match the user input can be created, and knowing what tids should be expected from a specific key search, the search operation can be run and seen if it returns the same results.

To test range search, enter "6" as the operation input and then the program will ask for the primary and secondary keys twice. This is to generate the minimum and maximum boundary keys for the range search. There is no error checking for logic or programmatically correct inputs. Once the inputs are given, the program will search through the tree and display all valid keys with their values. To test that it is correct, a print of the tree can be made and a manual check of which keys meet the range can be compared to the one generated by the program.

To quit the program, enter "7" as the operation input and the program will stop.

Example Test Cases

CASE 1 (loading tuples 1-8 and printing)

```
python Bplustree.py
```

```
1
```

```
netflixEdited.csv
```

```
1
```

```
8
```

```
2
```

CASE 2 (inserting and deleting)

```
python Bplustree.py
```

```
1
```

```
netflixEdited.csv
```

```
1
```

```
7
```

```
//inserting
```

3

8

3

9

3

10

//deleting

4

1

4

2

4

3